

Federated Visual Classification with Real-World Data Distribution

Tzu-Ming Harry Hsu, Hang Qi, Matthew Brown



ECCV 2020 #994

CIFAR 10/100

synthesized clients

low resolution

<1000 classes

EMNIST

natural clients

low resolution

<1000 classes

CelebA

not actual clients

large scale

high resolution

<1000 classes

iNaturalist-User-120k & iNaturalist-Geo

120k images

9275 clients

1203 classes

natural clients

large scale

high resolution

>1000 classes



Landmarks-User-160k

160k images

1262 clients

2028 classes

natural clients



large scale

high resolution

>1000 classes



FedVC: Splitting Imbalanced Clients with Virtual Clients




FedIR: Importance Reweighted Client Objectives

$$\mathbb{E}_p [\ell(x, y)] = \mathbb{E}_{q_k} \left[\ell(x, y) \frac{p(x, y)}{q_k(x, y)} \right]$$

target distribution

client's distribution



Thank You

tinyurl.com/federated-vision-datasets